




TESSERACT · SIH26003

Frontend design specification

Every screen of the AI cognitive gaming and memory assistance platform for elderly people living with dementia — patient mode and caregiver mode, with the interaction rules, states and data contracts each screen is built against.

SCREENS	MODES	VERSION	DATE
21	2	v0.1	5 Sep 2026

Design proposal for team review — not an approved specification.
No screen in this document has been built. Confirm against
PS003_MASTER_CONTEXT.md before coding.



Read the principles first. Every screen decision downstream is one of them applied.

01	Design principles	Eight rules the patient interface is built on, and why each one exists
02	Colour	Two mode palettes, contrast ratios, and what colour is never used to mean
03	Type & touch	Patient and caregiver scales, minimum targets, spacing
04	Components	Buttons, cards, chips, tiles, banners — the shared vocabulary
05	Two modes	Information architecture and the single door between them
06	Screen index	All 21 screens at a glance, with build ownership
07	Patient screens	P1–P8, one page each with callouts, states and data
08	Caregiver screens	C1–C13, same format
09	Flows	The demo golden path and the caregiver decision loop
10	States	Loading, empty, error, offline and insufficient data across the app
11	Build handoff	Screen → owner → work package → gate

What this document is not

It is not a clinical specification. Nothing here has been validated with patients, and no screen measures cognition. The metrics shown to caregivers are counts of how somebody played a game.

It is not a frozen contract either. Interaction rules and copy are proposals for the team to accept, change or reject at kickoff — the same way a recommendation works inside the product.

Conventions

Every measurement inside a phone frame is real Android **dp**. Frames are 360 × 740 dp, printed at 82 %.

Numbered dots on a mockup match the numbered notes beside it.

Conditional marks a screen that ships only if an earlier gate passes.

Sample content uses one synthetic family throughout: **Kamala** (72, the patient) and her daughter **Bidisha** (the caregiver). No real patient data appears anywhere in this document.

Eight rules. Every screen in this document is one of them applied.

These are design constraints, not aspirations. If a screen breaks one of them it is wrong, however good it looks — and the reviewer named on that screen's page is the person who should say so.

1

One decision per screen

A patient screen offers exactly one obvious action. Secondary routes exist but are visibly quieter — outlined, never filled. Two equally weighted buttons is a choice the patient has to hold in working memory.

2

Nothing is timed, ever

No countdown, no speed bonus, no session limit. Speed alone never raises difficulty. A pause is unlimited and a session can be left at any point without penalty.

3

Removal reads as loss

Matched pairs stay on screen, marked. Found words stay listed, struck through. Anything that disappears when the patient succeeds turns achievement into subtraction.

4

Help is furniture, not a confession

The Help control is permanent, full-size and introduced before it is needed. Using it flags the session assisted for analytics — and is never rendered to the patient or caregiver as a deficit.

5

No score, no streak, no comparison

The end screen reports one neutral count and one warm personal detail. There is no accuracy figure, no “better than last time”, and no reason to play again tonight.

6

Two channels for every recognition

A face carries a name. A place carries a label. Recognition is never made to depend on the image alone, because vision and familiarity vary independently of memory.

7

The caregiver decides, the app proposes

Every adaptive change is a proposal with its reasoning and source sessions attached. The patient's configuration does not move until a named human accepts it.






8

Say what the data is, where the data is

Screens that show numbers also say what those numbers are not. The interpretive boundary lives on the data screen, not in a policy page nobody opens.

Warm, high contrast, and never load-bearing.

Colour perception in the blue end of the spectrum declines with age and with several dementias, so the palette is warm-anchored and no meaning is carried by hue alone. Every state that uses colour also uses a word, a mark or a position.

 Ground #FDF8F1 Patient mode background. Warm off-white; never pure white, which glares under the bright screens older users tend to set.	 Surface #FFFFFF Cards and fields in both modes.	 Ink #1C1917 Primary text. 15.8:1 on ground.	 Ink muted #57534E Secondary text. 7.6:1 on ground. Never used for anything the patient must read.	 Line #E7E0D5 Hairlines, field borders, card outlines.	 Care primary #14594E Caregiver actions, caregiver-mode accents. 8.4:1 on white.
 Care tint #EAF3F0 Recommendation cards and caregiver highlights.	 Patient act #A8420B The single patient call to action, and the hand-over button. 6.9:1 on white.	 Patient tint #FBEDE0 Today's activity card, attention banners.	 Steady #3F7D5F Matched pairs, completion mark, saved state. Never "correct".	 Steady tint #E7F0EA Backgrounds for the above.	 Destructive #9B2C15 Delete only. Appears exactly once in the product.
Colour never means, on its own Right or wrong · better or worse · urgency · a clinical state. A mismatch in Reveal Match produces no red and no sound; the cards simply turn back.		The two-mode signal Teal is caregiver, warm orange is patient. The hand-over button on the caregiver home is orange — the one place the two meet — so the mode switch is legible before it is read.		Contrast floor 4.5:1 for all text and 3:1 for interactive boundaries, verified on the target device at maximum brightness. WCAG is the reference we design to; it is not a certification we claim.	

The patient scale is not the caregiver scale.

Two type systems in one app. Patient mode runs roughly 130 % of caregiver mode and honours the OS text-scale setting on top of that, so every patient layout must survive 200 % scaling without clipping.

PATIENT MODE — INTER

DISPLAY	30–34 sp / 700	Hello, Kamala
BODY	19–20 sp / 400	You can stop whenever you like.
BUTTON	22 sp / 700	Start
META	16–17 sp / 600	2 of 5

CAREGIVER MODE — INTER

TITLE	22 sp / 700	Kamala
BODY	13–14 sp / 400	Last synced 12 minutes ago
BUTTON	15 sp / 700	Yes, make this change
LABEL	10–11 sp / 700	LAST SESSION

Touch targets

Patient primary	68 dp tall, full width less 24 dp margins
Patient tile	118 dp square, 14 dp gutter
Patient chrome	44 dp minimum (Help, Pause)
Caregiver primary	52 dp
Caregiver chip / row	40–44 dp
Absolute floor	48 dp for anything a patient taps

Spacing

Patient screens: 24 dp side margin, 16–26 dp between blocks. Caregiver screens: 20 dp side margin, 12–16 dp between blocks. Two values each — no improvised gaps.


The scaling test

Every patient screen must be checked at the OS maximum text size on the target phone before its work package is called done. This is on Kovid's manual checklist, not a designer's discretion.


One vocabulary across both modes and every game.

Reveal Match and Word Find share their chrome exactly — the same Help pill, the same Pause, the same progress phrasing — so learning one game teaches the other.


PATIENT CONTROLS



CAREGIVER CONTROLS



BANNERS — EACH CARRIES A WORD, NOT JUST A COLOUR



One app, two experiences, one door between them.

Patient mode has no tab bar, no back-out gesture and no account. Caregiver mode holds everything else. The door is C10, and it is enforced by server-side access checks, not only by client navigation.

PATIENT MODE — NO NAVIGATION CHROME

P1 Patient home

The only entry point. One activity, one alternative.

P2 Choose an activity

At most three cards, all caregiver-approved.

P3 How to play

First session only; flagged as the tutorial.

P4 Reveal Match

The core loop. Every interaction emits an event.

P5 Paused

Unlimited. Resume, finish, or re-read the rules.

P6 Finished

One count, one warm detail, rest offered first.

P7 Word Find

Conditional — ships only after the Sep 7 gate.

P8 Rest state

Idle, opted out, or nothing approved.

C10 Hand over

The single door.
4 corner taps
+ caregiver PIN



Server enforces
the boundary too

CAREGIVER MODE — FOUR TABS

C1 Sign in

The only auth surface in the product.

C2 Add person — basics

Name, age, language, text size.

C3 Clinical context

Optional. Never sets difficulty.

C4 Know Me — people

Becomes game content.

C5 Know Me — words

15–20 target, skippable.

C11 Reminders • tab 3

Conditional. Local schedule.

C6 Home • tab 1

Pending suggestion, last session, sync, hand over.

C7 Sessions • tab 2

Comparable sessions only, exclusions shown.

C13 Insufficient data

Pre-threshold form of C7.

C8 Suggestion

Proposal, reasoning, three outcomes.

C9 Change it myself

Real override, safe range.

C12 Settings • tab 4


Queue, build ID, deletion.

Enrolment is a linear 4-step flow

C2 → C3 → C4 → C5, with a visible step counter and only one required field in the whole sequence. Target is under ten minutes; that target is a design goal, not a measured result.


Twenty-one screens, and who builds each.

Eight patient screens, thirteen caregiver screens. Two are conditional and ship only if the Sep 7 gate passes. Owner and work package are shown under each thumbnail; the full mapping is on the last page.



Patient home

The only screen a patient ever starts on. One greeting, one thing to do, one escape hatch. Nothing here can be mistapped into a caregiver area.



WHY IT LOOKS LIKE THIS

- 1 **Time-of-day greeting, not a date.** “Wednesday morning” orients without demanding recall of the date.
- 2 **Name from the profile, never a nickname the app invented.** Caregiver-entered display name only.
- 3 **One primary activity, pre-selected by the approved configuration.** The patient never chooses a difficulty.
- 4 **Duration and a permission to stop, stated before starting.** Removes time pressure by design.
- 5 **68 dp primary button, 22 sp label.** Single tap target, no competing CTA of equal weight.
- 6 **Secondary route is visibly quieter** — outlined, not filled, so there is always one obvious action.
- 7 **Caregiver attribution, not a login.** Reassures the patient someone they know arranged this.

INTERACTIONS

- Tap Start → P3 (how to play) on first ever session, P4 directly afterwards.
- Tap Something else → P2.
- No back gesture out of patient mode; exit requires the caregiver PIN (C10).

STATES


- **No approved activity yet** → P8 rest state.
- **Offline** → screen is unchanged; the last approved activity is served from cache.
- **Sync pending** → no indicator shown to the patient; it appears in caregiver mode only.

DATA & CONTRACT

- GET /v1/patients/{id}/activity → approved configuration or safe default.
- Served from the local cache when offline; never blocks on the network.

Choose an activity

A short, closed list. At most two or three activities are ever shown, because a long menu is a memory task in itself.



WHY IT LOOKS LIKE THIS

- 1 **Question as the heading, in the patient's second person.** No app jargon, no "Games".
- 2 **Maximum three cards.** Anything beyond the approved set is hidden, not greyed out.
- 3 **Each card carries a plain-language subtitle** that says what it contains, not how hard it is.
- 4 **Image chip is drawn from the patient's own Know Me content** so the card is recognisable before it is read.
- 5 **Word Find only appears once S19 ships** and only if the patient has 8 or more usable words.
- 6 **Opting out is a first-class action**, never a dead end or a nag.

INTERACTIONS

- Tap a card → P3 on first play, otherwise P4.
- Tap Not right now → P8, which does not scold or re-prompt.

STATES


- **Only one activity approved** → this screen is skipped entirely; P1 goes straight to the game.
- **Word Find unavailable** (too few words, or S19 not shipped) → card is absent, not disabled.

DATA & CONTRACT

- Same activity payload as P1; the list is whatever the caregiver has approved.

How to play

Shown before the first ever session of a game, and re-shown on request. This session is flagged as the tutorial and is excluded from the baseline.



WHY IT LOOKS LIKE THIS

- 1 **Three steps, never more.** Each is one sentence in the imperative.
- 2 **Numerals in filled circles**, 40 dp, so the sequence reads without depending on line order.
- 3 **"Nothing is timed" is stated explicitly**, because the expectation of timing is itself the stressor.
- 4 **Help is introduced before it is needed**, so requesting it never feels like failure.
- 5 **Begin, not Start** — distinct word from P1 so the two screens are not confusable.

INTERACTIONS

- Tap Begin → P4, with `is_tutorial: true` on the session.
- Reachable again from P5 (paused) via "How do I play?".

STATES


- **Repeat sessions** skip this screen; it is never forced twice.
- **Caregiver-triggered replay** from C12 resets the tutorial flag.

DATA & CONTRACT

- Session created with `is_tutorial: true`; analytics must exclude tutorial sessions from the provisional baseline (S17).

Reveal Match — in play

The screen the entire product is measured on. Every interaction here emits an event; nothing here scores, ranks or corrects the patient.



WHY IT LOOKS LIKE THIS

- Help is a permanent, full-size button**, not a hidden menu item. Emits `hint_requested` and flags the session assisted.
- Pause is always one tap away** and is the only other chrome on the screen.
- Instruction restated on the play screen** so it survives a lapse mid-session.
- Matched pairs stay visible and marked** rather than disappearing — removal reads as loss.
- Face-up cards are labelled with the person or place name**, so recognition is supported by two channels.
- 118 dp tiles, 14 dp gutters, 3 columns at level 1**. Level only changes grid size, never speed.
- Progress is stated as pairs found, never as a score or accuracy**. No streaks, no stars, no timer.

INTERACTIONS

- Tap card → `card_revealed`. Second card resolves → `pair_resolved` with `matched true/false`.
- Mismatch: both cards turn back after ~1.2 s with no sound or colour of failure.
- Help → reveals one true pair, emits `hint_requested`, sets `assisted: true` for the session.
- Pause → P5, emits `paused`; resuming emits `resumed`.

STATES

- Backgrounded** → treated as a pause; that interval is excluded from active duration.
- Offline** → identical behaviour; events queue to SQLite before any upload attempt.
- Killed mid-session** → session status interrupted; it is excluded from the baseline.

DATA & CONTRACT

- Events carry `event_id`, `seq`, `elapsed_ms`, `difficulty`, `input_mode`, `assisted`.
- Never place patient names, personal words or photo URLs in the telemetry payload.

Owner: Shanks


Reviewer: Pranav

S11

Core · Sep 6 gate

Paused

A pause must never feel like an interruption or a penalty. This screen exists so that stopping is a supported, dignified action.



WHY IT LOOKS LIKE THIS

- 1 **"Taking a break", not "Paused"**. Describes the person, not the software state.
- 2 **Explicit reassurance that nothing is lost**. This is the single most common anxiety at this moment.
- 3 **Resume is the primary action**, but finishing is offered at equal size and without warning language.
- 4 **Finishing early is a valid completion**: the session is stored with status ended_early and is still shown to the caregiver.
- 5 **Route back into instructions**, because a pause is often caused by losing the thread.

INTERACTIONS

- Carry on → P4, emits resumed; the paused interval is subtracted from active duration.
- Finish for now → P6, session completed with the events collected so far.
- How do I play? → P3, which does not restart the session.

STATES

- Auto-pause on backgrounding** lands on this exact screen, so returning is never disorienting.
- Long pause (> 10 min)** is still resumable; the app makes no decision to end the session on the patient's behalf.

DATA & CONTRACT

- paused / resumed events with elapsed_ms; active duration excludes every paused interval (S13).

Owner: Shanks


Reviewer: Pranav

S11

Core

Session finished

The end screen carries no evaluation. It reports one neutral count and one warm personal detail, then offers rest as the primary action.



WHY IT LOOKS LIKE THIS

- 1 **A completion mark, not a grade.** Same mark whether five pairs were found or one.
- 2 **"All done", never "Well done" or "You scored".** Praise implies a threshold was met.
- 3 **One count only — pairs found.** No accuracy, no time, no comparison to last time.
- 4 **Names the caregiver,** closing the loop the patient was told about on P1.
- 5 **A personal, non-evaluative observation** drawn from the session. Warmth without measurement.
- 6 **Rest is the primary action.** Playing again is available but deliberately secondary — no "keep the streak" pressure.

INTERACTIONS

- Rest now → P8. Play once more → a new session (still counted separately).
- Nothing on this screen leads into caregiver areas.

STATES

- **Ended early** → identical layout, count reflects what was played, wording unchanged.
- **Zero pairs found** → the count line is replaced with "You had a go today, Kamala." Never "0 pairs".
- **Offline** → identical; the "Bidisha will see" promise is honoured on next sync.

DATA & CONTRACT

- POST /v1/sessions/{id}/complete with final_seq and status.
- The screen renders from local state, never from a server round-trip.

Owner: Maharshitha


Reviewer: Shanks

S11

Core

Personalized Word Search

The clearest visible proof of personalisation: the grid is built from this patient's own meaningful words. Ships only after the full loop passes on Sep 7.



WHY IT LOOKS LIKE THIS

- 1 The word list is the patient's own **Know Me words** — family names and places, never a generic dictionary.
- 2 Found words stay on screen, **struck through**, so progress is visible without a score.
- 3 Four words maximum per round at level 1; the grid grows before the list does.
- 4 Uppercase, 19 sp, 6×6 grid. Letter cells are ~48 dp so a drag is forgiving.
- 5 Found runs are **highlighted, not removed**. Same visual language as matched pairs in P4.
- 6 Same Help and Pause chrome as P4 — one interaction vocabulary across every game.
- 7 Progress phrased **identically to P4** ("1 of 4 found"), so the two games feel like one product.

INTERACTIONS

- Drag across letters to select a run; a correct run locks and strikes the chip.
- Help → reveals the first letter of one remaining word; sets `assisted: true`.
- Events reach the identical session pipeline as P4 — no separate analytics path.

DATA & CONTRACT


- Personal words are read locally; the word text must not appear in telemetry payloads.
- Same sessions / events contract as Reveal Match, distinguished by `game_id`.

STATES

- **Fewer than 8 usable words** → the game is not offered at all (card absent from P2).
- **Words too long for the grid** → excluded at generation time; never truncated on screen.
- **Duplicate or single-character words** → filtered during content generation.

Rest state

Where the patient lands after finishing, after declining, or when no activity is approved. It must never read as an error or an empty container.



WHY IT LOOKS LIKE THIS

- 1 **A calm mark, not an empty-state illustration** and never a sad face or a warning triangle.
- 2 **States the situation without implying anything is wrong** with the patient or the app.
- 3 **Names who will act next.** Removes the sense of being stranded.
- 4 **Handing the phone over is the exit**, phrased as an action the patient takes, not a locked door.

INTERACTIONS

- Show Bidisha this phone → caregiver PIN prompt → C6.
- Pull to refresh is deliberately absent; the patient is never asked to troubleshoot sync.

STATES

- This screen is itself the empty state for patient mode — there is no separate error screen shown to the patient.
- **Server unreachable** → identical screen; failure is never surfaced to the patient.

DATA & CONTRACT

- Rendered when activity is null or the patient has opted out of the current session.

Owner: Maharshitha

Reviewer: Shanks

S10

Core

Caregiver sign in

The only authentication surface in the product. Establishes from the first screen that the patient is never asked to hold credentials.

WHY IT LOOKS LIKE THIS

- 1 **Mark and product name, no marketing claim.** Nothing here should assert clinical benefit.
- 2 **Says who this screen is for,** in the first sentence, to prevent a patient attempting it.
- 3 **Mobile number as the primary identifier** — more reliable than email for the target caregivers.
- 4 **Standard 52 dp caregiver control height;** caregiver mode uses normal mobile density, not the patient scale.
- 5 **The patient/caregiver boundary is stated on the sign-in screen itself,** not buried in settings.

INTERACTIONS

- Sign in → C6 if a patient exists, otherwise C2.
- Create an account → same fields plus a confirm step, then C2.

STATES

- **Wrong credentials** → inline message under the field, never a modal.
- **Offline** → first sign-in requires connectivity and says so plainly; an already-enrolled device does not.
- **Token expired** → returns here, and patient mode continues to work from cache.

DATA & CONTRACT

- Firebase Authentication issues the identity; FastAPI verifies the token and the caregiver–patient membership on every route.
- 401 identity required, 403 no access to this patient.

Owner: Pranav

Reviewer: Shanks

S08

Core

Add a person — basics

The shortest possible path to a usable profile. Only the display name is genuinely required; everything else can be changed later.

9:41

< Add a person 1 1 of 4

This is the person who will use the games.

What should we call them?

Kamala

This name is what appears on their screen.

Age

72

Language on their screen

English Assamese Bengali

Assamese and Bengali packs are in review and are marked in the app until a native reviewer signs them off.

Text size

Large Extra large

Continue

WHY IT LOOKS LIKE THIS

- Four-step counter set from the start**, so the caregiver can see how long enrolment is. Target is under ten minutes end to end.
- "What should we call them?"** — invites the familiar name the patient responds to, not a legal name.
- Explains where the value will appear**, so the caregiver chooses deliberately.
- Language pack honesty**. Unreviewed packs are labelled here and in the app; we do not claim language support we have not had reviewed.
- Text size is a profile setting, not a device setting**, and defaults to Extra large.

INTERACTIONS

- Continue → C3. Back preserves entered values.
- Only the name blocks Continue; age and language have safe defaults.

STATES

- Duplicate name under the same caregiver** → allowed, disambiguated by age in lists.
- Offline** → the whole flow is completable offline and syncs on reconnect.

DATA & CONTRACT

- POST /v1/patients → patient_id and version.
- Fields: display_name, age, language, accessibility.text_scale.

Clinical context — optional

Records what a clinician has already said. It is stored as context only: it never sets difficulty, and nothing in the product infers it from gameplay.

9:41

< Add a person 2 of 4

Anything a doctor has told you?

1 You can skip all of this. It does not change the games.

2

Type, if you know it

3 Not known Alzheimer's

4 Vascular Lewy body

Frontotemporal Mixed / other

5 Stage, if you know it

Not known Early Middle

Later

6

Tesseract does not diagnose, and does not work out a diagnosis from how someone plays. What you enter here is kept as context for you and any health worker you share it with.

7 Continue

Skip this

WHY IT LOOKS LIKE THIS

- 1 **"Anything a doctor has told you?"** frames this as reporting, not self-assessment.
- 2 **Skippability is stated before the first control**, not discovered at the bottom.
- 3 **"Not known" is pre-selected and sits first**, at equal visual weight — the honest default for most families.
- 4 **All five recognised types are offered**, plus mixed/other, so the caregiver is never forced into a wrong bucket.
- 5 **Stage is equally optional** and equally uncoupled from difficulty.
- 6 **The no-diagnosis boundary appears at the point of data entry**, where it actually matters, not only in a policy screen.
- 7 **Skip is a plain visible action**, never a greyed link.

INTERACTIONS

- Continue or Skip this → both go to C4; skipping stores explicit nulls, not empty strings.
- Values are editable at any time from C12.

STATES

- **All fields left unknown** is the expected common case and must be fully supported downstream.
- No screen anywhere in the product may render a stage the caregiver did not enter.

DATA & CONTRACT

- known_type, known_stage — nullable, free of inference.
- Recommendation rules (S16) must not read these fields. Enforce this in code review.

Owner: Maharshitha


Reviewer: Pranav

S09 / S10

Core

Know Me — people and places

Where the personalisation actually comes from. Everything added here becomes game content, so the screen states that consequence up front.



WHY IT LOOKS LIKE THIS

- Says exactly what the input is used for** — the caregiver is choosing game content, not filling a form.
- Counts in the section label**, so progress is visible without a progress bar.
- Relationship shown under every face**. It is used in the game caption and in P6's personal line.
- Photos are optional** — a coloured initial is a complete, dignified fallback, not a broken image.
- Places are first-class content**, not an afterthought, because visuospatial recall is a target area.
- One combined add action**, so the caregiver is not asked to categorise before they start.

INTERACTIONS

- Add → a sheet with name, relationship, optional photo from camera or gallery.
- Long-press an entry to edit or remove it.
- Continue is enabled with zero entries; the games fall back to bundled neutral imagery.

STATES

- No photos added** → initials avatars everywhere; the product must be fully demonstrable without a single photo.
- Photo too large** → downscaled on device before storage.
- Deleted person** → removed from future game content; past sessions keep their integer counts, not the image.

DATA & CONTRACT

- PUT /v1/patients/{id}/personalization.
- Photos are stored locally first; treat them as personal data — never in telemetry, never in an AI prompt.

Owner: Maharshitha


Reviewer: Shanks

S09 / S10

Core

Know Me — meaningful words

Collects the 15–20 word target without making it a requirement. The counter encourages; it never blocks or scolds.



WHY IT LOOKS LIKE THIS

- 1 **Single-field entry with an inline Add** — the fastest possible repeated input on a phone.
- 2 **Chips are individually removable**, so a mistyped word costs one tap to fix.
- 3 **Progress toward 15–20 is shown as a soft bar**, never as “10/20 required”.
- 4 **The copy explicitly says fewer is fine**. Ten meaningful words beat twenty forced ones.
- 5 **Names from C4 are folded in automatically**, so the caregiver never types the same name twice.
- 6 **“Do this later” is offered at the final step**; setup must be completable without any words at all.

INTERACTIONS

- Add → appends a chip and clears the field, keeping focus for fast repeated entry.
- Finish setup → C6. Do this later → C6 with a gentle prompt on the patient card.

STATES


- **Zero words** → Word Find is simply not offered; Reveal Match is unaffected.
- **Duplicate word** → silently ignored, chip flashes rather than an error appearing.
- **Word under three letters or over ten** → accepted for storage, excluded from grid generation with an inline note.

DATA & CONTRACT

- `personal_words(id, patient_id, text, locale)`.
- Word text is patient data: it must not appear in telemetry or in any AI prompt (S19).

Caregiver home

Answers three questions in one glance: does anything need me, how did the last session go, and is the data safe. Then hands the phone over.



WHY IT LOOKS LIKE THIS

- 1 **Patient name is the page title**; the caregiver's own identity is secondary.
- 2 **A pending recommendation is the top item on the screen**, because human review is the product's core claim.
- 3 **Last session in plain counts**: pairs, time, help. No accuracy percentage, no invented domain score.
- 4 **"Help used: Not used" is stated positively** — using help is never rendered as a deficit.
- 5 **Sync state is explicit and always visible**: what is saved and when it last reached the server.
- 6 **Handing over is the primary action of the whole screen**, in the patient's warm accent so the two modes are colour-coded.
- 7 **Four-tab navigation exists only in caregiver mode**. Patient mode has no tab bar at all.

INTERACTIONS

- Review → C8. See all sessions → C7. Hand the phone over → C10.
- Tapping the sync card → C12.

STATES

- **No pending recommendation** → the amber banner is absent, not replaced by an empty placeholder.
- **Fewer than three sessions** → the last-session card renders, but C7 shows the insufficient-data state (C13).
- **Never synced** → card reads "Saved on this phone, not sent yet" with the queue count.

DATA & CONTRACT

- GET /v1/patients/{id}/summary and /recommendations.
- Renders from cache first, then reconciles; it must never show a spinner as the whole screen.

Owner: Maharshitha


Reviewer: Shanks

S15

Core

Session summary

Shows observable game performance over time, with the comparability rules and the exclusions visible. This is where the product's honesty is most tested.



WHY IT LOOKS LIKE THIS

- 1 **Only comparable sessions are charted** — same game, same version, same level, same input mode.
- 2 **The axis is labelled in the caregiver's language**, not as "accuracy".
- 3 **Sessions counted vs left out are shown side by side**. Hiding exclusions would misrepresent the trend.
- 4 **Assisted sessions are counted separately**, never silently blended into the same line.
- 5 **The reason for each exclusion is written out**, so an absent bar is never mysterious.
- 6 **The interpretive boundary is on the data screen itself**. This sentence is a release blocker if removed.

INTERACTIONS

- Tap a bar → that single session's detail: events summary, configuration, timestamps.
- Filter by game when more than one game has data.

STATES

- **Fewer than three comparable sessions** → C13 insufficient-data state replaces the chart.
- **Configuration changed** → a new comparison series starts; the chart shows the break rather than joining across it.
- **Seeded demo data** → must carry a visible "demonstration data" label.

DATA & CONTRACT

- GET /v1/patients/{id}/summary returns metrics, sample counts, comparability filters and last sync.
- Metrics come from S13; zero attempts yields null, never a zero.

Owner: Ruthika


Reviewer: Pranav

S13 / S15

Core

Recommendation review

The screen that makes the product defensible. A rule proposes, a human decides, and the reasoning is shown in full before the decision is asked for.



WHY IT LOOKS LIKE THIS

- 1 **The proposal is stated as a concrete configuration change**, never as a judgement about the person.
- 2 **"Why" is written in the caregiver's language**, generated from the stored rule inputs — not a raw threshold dump.
- 3 **"Based on" names the exact source sessions**, so the caregiver can go and check them.
- 4 **Rule version is recorded and displayed**. Every stored decision is reproducible after the fact.
- 5 **Three outcomes, all first-class**: accept, modify, reject. Rejecting is not styled as the wrong answer.
- 6 **The status quo while pending is stated explicitly** — the patient is never silently moved.

INTERACTIONS

- Yes → the new configuration becomes the approved activity, effective next session.
- Change it myself → C9. Leave it as it is → proposal stored as rejected, with the caregiver's identity.
- Difficulty never changes mid-session, under any outcome.

STATES

- **Insufficient data** → no proposal is generated at all; the banner does not appear on C6.
- **Assisted or interrupted sessions in the window** → automatic progression is blocked by rule.
- **A second proposal arriving while one is pending** → supersedes it; only one is ever shown.

DATA & CONTRACT

- GET /v1/patients/{id}/recommendations, POST /v1/recommendations/{id}/decision.
- Stored: rule_version, source_session_ids, observed_values, reason, status, decided_by.

Owner: Pranav

Reviewer: Shanks

S16 / S18

Core · Sep 7 gate

Modify the recommendation

Makes caregiver override real rather than rhetorical. Every value the rule could set is directly editable here, within a safe allowed range.

9:41

< **Change it myself**

You decide what Kamala plays next. Nothing here is locked.

1

How many pairs

2 3 4 5 6 8

Suggested: 6

3 Extra support

☒ Show names under the pictures

4 ☐ Keep matched pairs turned over

Which game

Picture Match Word Find

5

Save for Kamala

6 Takes effect the next time she plays.

WHY IT LOOKS LIKE THIS

- 1 **Opens by affirming the caregiver's authority.** This is the sentence that makes the override credible.
- 2 **Discrete levels, not a slider.** A slider implies a continuous scale of ability, which we do not claim to measure.
- 3 **The rule's suggestion stays visible but unselected,** so the caregiver sees what they are departing from.
- 4 **Support options are separate from difficulty.** Adding help is not the same as making a game easier.
- 5 **Only games the patient has content for are offered;** Word Find is disabled without enough words.
- 6 **The effective time is stated:** never mid-session, always from the next session.

INTERACTIONS

- Save → stores a modify decision with the caregiver's chosen configuration and returns to C6.
- Values outside the allowed range are not offered; the UI cannot produce an invalid configuration.

STATES

- **Reached without a pending proposal** (from C12) → the same screen with no "Suggested" hint.
- **Offline** → the decision is queued and applied locally at once, so the patient is not blocked.

DATA & CONTRACT

- POST /v1/recommendations/{id}/decision with modified_config.
- Server re-validates the range; a client-only guard is not sufficient.

Owner: Pranav


Reviewer: Shanks

S18

Core

Hand over to patient mode

A deliberate, explained transition. It is the only door between the two modes, and it tells the caregiver exactly what will be hidden and how to return.



WHY IT LOOKS LIKE THIS

- 1 **Full-colour interruption screen.** A mode change is significant enough to deserve the whole screen.
- 2 **States precisely what the patient can and cannot reach.** No vague “restricted mode” language.
- 3 **The return gesture is taught before it is needed,** not left to be rediscovered under stress.
- 4 **Four corner taps plus a PIN** — deliberately hard to trigger accidentally, easy for the caregiver to repeat.
- 5 **Backing out is offered;** the caregiver is never trapped in the handover.

INTERACTIONS

- Hand it over → P1, with caregiver navigation removed from the stack, not merely hidden.
- Return path → PIN sheet → C6. Three wrong PINs → sign in again with the account password.

STATES

- **PIN not yet set** → this screen collects one before proceeding.
- **Patient mode across a restart** → the app relaunches into patient mode; the state is persisted, not held in memory.

DATA & CONTRACT

- Patient mode must be enforced by server-side access checks too, not only by client navigation.
- A patient-mode device session carries only the approved-activity scope.

Owner: Shanks


Reviewer: Pranav

S18 / S22

Core

Routine reminders

A minimal locally scheduled reminder list. It ships only if the core loop is stable, and it is explicit about what an acknowledgement does and does not prove.



WHY IT LOOKS LIKE THIS

- 1 **Caregiver-written routine text**, kept deliberately distinct from anything resembling medical advice.
- 2 **Toggle rather than delete** for the common case of pausing a routine.
- 3 **The adherence boundary is stated on the screen that creates the expectation.** This wording is a claim-register item.
- 4 **Local scheduling is explained** so the caregiver understands why it survives being offline.
- 5 **Permission denial is named as a real outcome**, not discovered when a reminder silently fails.

INTERACTIONS

- Add → label, time, repeat days. No medication database, no dosage fields.
- The patient sees a full-screen reminder with a single Done action.

STATES

- **Notification permission denied** → a persistent inline warning with a route to Android settings.
- **Device restarted** → reminders must be rescheduled on boot; this is an explicit test case.
- **Time zone changed** → reminders keep local wall-clock time.

DATA & CONTRACT

- Stored locally first; server copy is for caregiver continuity across devices only.
- Acknowledgement is stored as `reminder_seen`, never as `medication_taken`.

Owner: Maharshitha


Reviewer: Shanks

S20

Conditional

Settings, sync and data

Everything reversible in one place: the queue state, the profile content, and deletion. Build identity is visible here so QA can report against it.



WHY IT LOOKS LIKE THIS

- 1 **Queue depth and last successful sync as plain counts** — the caregiver's only window into offline state.
- 2 **App version and build number are visible**, because Kovid's QA checklist requires a build ID with every report.
- 3 **Every onboarding step is re-enterable**; nothing collected during setup is permanently frozen.
- 4 **Replaying the practice round resets the tutorial flag** and starts a fresh comparison series.
- 5 **Where data lives is stated in ordinary words**, distinguishing on-device content from server records.
- 6 **Deletion is a real, reachable action**, not a support-email request. Tombstoned so a reconnect cannot restore it.

INTERACTIONS

- Delete → typed confirmation naming the patient, then a summary of what will be removed from device and server.
- Changing the activity here creates a caregiver decision record just as C9 does.

STATES

- **Queue non-empty** → the count is shown on C6 as well, so it is never only discoverable here.
- **Deletion while offline** → applied locally, queued as a tombstone, reconciled on reconnect.
- **Partial deletion failure** → reported honestly with what remains, never silently reported as complete.

DATA & CONTRACT

- Deletion must cover server records, photos, cached copies and the outbox.
- Full deletion propagation stays labelled *planned* until it is implemented and tested (S22).

Owner: Pranav


Reviewer: Shanks

S22

Core

Insufficient data

The default state of the whole product for its first days of use. Getting this right matters more than the chart, because most demos and most real first weeks land here.



WHY IT LOOKS LIKE THIS

- 1 **A neutral mark, never a warning icon.** Having little data is normal, not a problem to fix.
- 2 **Names the actual condition** — three comparable sessions — rather than saying “no data”.
- 3 **Progress toward the threshold is visible**, so the caregiver knows the app is working and what to expect.
- 4 **Sessions already played are listed with their status**, including why the practice round does not count.
- 5 **The threshold is labelled as an engineering choice**, not a clinical protocol. This is a claim-register item.

INTERACTIONS

- No call to action pushing the patient to play more — the caregiver decides the pace.
- Tapping a listed session opens its detail, same as C7.

STATES

- This *is* a state screen; it is the pre-threshold form of C7.
- **Zero sessions** → same layout, list replaced by “Nothing played yet”.
- **Configuration changed after two sessions** → the count restarts and the screen says so.

DATA & CONTRACT

- Baseline uses the first three valid comparable sessions after the tutorial (S17).
- Recommendations are not generated at all in this state — C6 shows no banner.

Owner: Ruthika

Reviewer: Pranav


S17

Core


The demo golden path.

This is the sequence the Sep 8 build must repeat twice from clean synthetic data. Every box below is a screen in this document; if any one of them is missing, the demo is a mock.


SETUP — CAREGIVER, ROUGHLY SIX MINUTES



PLAY — PATIENT, ABOUT FOUR MINUTES



REVIEW — CAREGIVER, THE PART THAT WINS THE ROUND




What must be said out loud during this demo

That the historical chart is seeded demonstration data. That nothing here diagnoses. That the language packs are unreviewed. Aryan owns the script; the claim register (S01, Kovid) is what it is checked against.

The caregiver decision loop, in full.

The loop that distinguishes this product from a game catalogue. Read left to right; the two guard conditions in the middle are where most of the engineering honesty lives.



The states that decide whether this feels finished.

A hackathon build is usually judged on its happy path and lost on its empty one. Every screen owner must implement the row for their screen before the work package is called done.

STATE	PATIENT MODE	CAREGIVER MODE	RULE
Loading	Never a spinner covering the screen. P1 renders from cache instantly; the patient waits for nothing.	Cached content renders first, then reconciles. Skeleton blocks only inside cards, never full-page.	No patient screen may ever depend on a network round-trip to appear.
Empty	P8 rest state. Calm mark, plain sentence, names who acts next. Never an illustration of sadness.	C13 insufficient data. Names the actual threshold and shows progress toward it.	An empty state explains the condition and who resolves it. It never says “No data”.
Error	Not shown. Failures are absorbed silently and surfaced to the caregiver instead.	Inline under the relevant control, with a stable error code and a request ID for support.	The patient is never asked to troubleshoot. That is a hard rule, not a preference.
Offline	Identical behaviour. Games play, events queue to SQLite, the approved activity is served from cache.	C6 sync card reads “Saved on this phone, not sent yet” with the queue count. C12 shows the depth.	Offline is a normal operating mode, not a degraded one. It is never styled as a warning to the patient.
Interrupted	Relaunching returns to P5 paused, not to a lost session. Events already queued survive the restart.	The session appears in C7 as excluded, with “interrupted” stated as the reason.	Interrupted sessions are stored and shown, never deleted — but they never feed the baseline.
Insufficient	No effect. The patient's experience does not change with the amount of data collected.	C13 replaces the chart; no proposal is generated; no banner appears on C6.	The app holds rather than guesses. Silence is the correct output below threshold.
Seeded	Not applicable.	Any seeded historical trend carries a visible “demonstration data” label on the chart itself.	A release is blocked if seeded data can be mistaken for real data. Kovid's checklist covers this.
Permission denied	A reminder simply does not appear. No system dialogue is shown to the patient.	Persistent inline warning on C11 with a route into Android settings.	Denied permissions are surfaced where they can be fixed, by the person who can fix them.

Screen to owner to work package.

Ownership follows **PS003_ROADMAP.md**. Nothing here changes who owns what — it only says which screens fall inside each existing work package.

ID	SCREEN	OWNER	REVIEWER	PACKAGE	GATE
P1	Patient home	Maharshitha	Shanks	S10	Core
P2	Choose an activity	Maharshitha	Shanks	S10	Core
P3	How to play	Shanks	Pranav	S11	Core
P4	Reveal Match — in play	Shanks	Pranav	S11	Core · Sep 6 gate
P5	Paused	Shanks	Pranav	S11	Core
P6	Session finished	Maharshitha	Shanks	S11	Core
P7	Personalized Word Search	Shanks	Pranav	S19	Conditional · after Sep 7
P8	Rest state	Maharshitha	Shanks	S10	Core
C1	Caregiver sign in	Pranav	Shanks	S08	Core
C2	Add a person — basics	Maharshitha	Shanks	S09 / S10	Core
C3	Clinical context — optional	Maharshitha	Pranav	S09 / S10	Core
C4	Know Me — people and places	Maharshitha	Shanks	S09 / S10	Core
C5	Know Me — meaningful words	Maharshitha	Pranav	S09 / S10	Core
C6	Caregiver home	Maharshitha	Shanks	S15	Core
C7	Session summary	Ruthika	Pranav	S13 / S15	Core
C8	Recommendation review	Pranav	Shanks	S16 / S18	Core · Sep 7 gate
C9	Modify the recommendation	Pranav	Shanks	S18	Core
C10	Hand over to patient mode	Shanks	Pranav	S18 / S22	Core
C11	Routine reminders	Maharshitha	Shanks	S20	Conditional
C12	Settings, sync and data	Pranav	Shanks	S22	Core
C13	Insufficient data	Ruthika	Pranav	S17	Core

Before anyone opens an editor

Three decisions are still open and each one blocks parallel work on these screens: the Flutter state-management approach, the backend base URL the app will hit on Sep 6, and whether the v0.1 contracts exist as committed JSON fixtures rather than prose.

Until the fixtures are files in the repository, “contracts first” has been described but not done, and P1, C6 and C7 cannot be built against anything real.

Suggested build order for the frontend

Shell and one fixture-backed screen → P1 → P4 → P6 → C6, then C2–C5, then C7 and C8. P5, P8 and C13 are small and belong alongside their parents — they are the screens the demo actually lands on.

What would make this document wrong

A screen that scores the patient. A chart of a cognitive domain. A difficulty that changes without a caregiver decision. A language claimed as supported without a native reviewer. Any of these is a release blocker, not a design preference.

Status

All 21 screens are **NOT STARTED**. This is a design proposal for team review, produced 5 September 2026. No application code exists yet.