

# Tesseract / Additions only

Seven additions with full implementation detail / 6 September 2026

## How to read this edition

**Existing:** the planned nine games, P1-P9 patient screens, C1-C7 caregiver screens, D1-D8 doctor screens, required reminders, analytics and the 40-package roadmap. Existing does not mean implemented. The existing specification remains in the separate product/build documents; it is not reproduced here. This document retains the full reviewed detail for all seven additions.

**Added:** A1-A7 from PS003\_ADDITIONS.md, with corrections identified during review. The user requested their inclusion in a new PDF. Their detailed technical rules, owner handoffs and priority order below are recommendations; no code or clinical capability is established.

| Addition                 | Recommendation / relation to existing scope                                                                        |
|--------------------------|--------------------------------------------------------------------------------------------------------------------|
| A1 Conversation prompts  | Extend C4 with factual session observations. Start narrow; never infer liking or remembering.                      |
| A2 Played together       | Extend C5 and shared wording. Separate companion/support context in analytics.                                     |
| A3 Content curation      | Extend C3 and content selection with accurate usage counts and rotation.                                           |
| A4 Accessibility fitting | Extend tutorial/settings with optional reviewed interface suggestions. Measurement design must be corrected first. |
| A5 Reminder bridge       | Optional short activity after an acknowledged reminder; separate reminder and game records.                        |
| A6 Daypart preferences   | Optional caregiver schedule. Record context; initially annotate rather than split every trend.                     |
| A7 Appointment summary   | New caregiver preview with a deterministic one-page PDF export. Separate from D7.                                  |

## What does not change

Nine games remain allocated Aryan five, Ruthika two and Pranav two using the working mapping. Shanks retains shared frontend architecture/integration; Pranav retains backend/AI architecture. No new patient screen is needed. Unmeasured free play remains outside the selected additions.

## Main review conclusion

The strongest early additions are A2 and a narrow A1. A7 is valuable once reliable summaries exist. A3 needs accurate content events, and A4 requires a prototype before a release commitment. The original hour/half-day estimates omit cross-game integration and QA; they are not delivery promises.

# Patient screens / existing and changed

The patient still has P1-P9

| Screen                   | Existing purpose                  | Addition / acceptance                                                                                            |
|--------------------------|-----------------------------------|------------------------------------------------------------------------------------------------------------------|
| P1 Home                  | Greeting and Start activity       | A6 selects an eligible preferred activity; retain offline fallback and patient choice.                           |
| P2 Choose Activity       | Small readable activity selection | Keep structure; availability and caregiver preferences must not create an empty forced path.                     |
| P3 How to Play           | Repeatable tutorial               | A2 companion wording. A4 may add an optional known-target tutorial interaction; it is a real interaction change. |
| P4 Game                  | Shared play, help, break, finish  | A2 wording; A5 short configuration from the shared host. Do not implement nine independent bridge variants.      |
| P5 Taking a Break        | Preserve paused session           | Preserve companion, origin and configuration on resume; no timer reset that corrupts metrics.                    |
| P6 Finished              | Gentle finish and exit            | A2 companion wording; A5 offers a calm return without additional activity pressure.                              |
| P7 Personalized Activity | Approved familiar suggestion      | A3 rotation and A6 preferences apply within approved content/support constraints.                                |
| P8 Rest                  | No active session                 | May receive return after a declined bridge when originally at rest. Never strand an existing paused game.        |
| P9 Simple Progress       | Gentle activity history           | Keep structure. Short and companion activities receive neutral acknowledgement, without rankings or penalties.   |

## Patient reminder state

A5 adds an optional follow-on card with a clear decline after acknowledgement, controlled by a caregiver setting that defaults off. It is a new UI state even though it does not require another numbered patient screen. Decline returns to the previous safe context, such as Home, Rest or an existing paused activity.

## Protect real-world routines

The activity offer should not distract from the action the reminder is prompting. Start with a caregiver-selected non-urgent use case or defer the offer. A game never blocks reminder delivery or acknowledgement, and acknowledgement never proves the real routine happened.

# Caregiver and doctor changes

Existing C1-C7 and D1-D8 remain; C8 is a proposed new preview

| Screen                 | Change                                                                                                                                                       |
|------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|
| C1 / C2                | Keep sign-in and basics. Accessibility preferences remain directly editable.                                                                                 |
| C3 Know Me             | A3 item presentation/choice/skip counts, content rotation context and edit/remove access.                                                                    |
| C4 Home / detail       | A1 conversation observation, companion labels, occasional content/accessibility suggestion and appointment-summary entry. Avoid stacking every card at once. |
| C5 Hand Over           | A2 played-together toggle and handover preview label.                                                                                                        |
| C6 Reminders           | A5 per-reminder short-activity offer, default off; show acknowledgement separately from bridge state.                                                        |
| C7 Settings            | A4 caregiver-set/accepted-suggestion history and reset; A6 optional daypart schedule. Prefer one canonical setting location.                                 |
| C8 Appointment Preview | NEW proposed screen: date range, field selection, exact PDF preview and explicit share/print action. Can alternatively be a C4 subview.                      |
| D1 / D2 / D8           | Keep identity, assigned caseload and notes structures. Existing permissions/audit remain.                                                                    |
| D3 / D5                | Include support/session-type context in overview and grouped observations. Do not silently mix modes in summaries.                                           |
| D4 / D6                | Label/filter companion and short activities; annotate daypart and configuration changes.                                                                     |
| D7 Report              | Use the same eligibility rules and labels as charts. It remains a doctor-side generated draft with review provenance.                                        |

## Screen count and design handoff

If C8 is a standalone screen, the total is 25: nine patient, eight caregiver and eight doctor screens. Maharshitha updates Figma and bounded components; Shanks owns navigation, shared strings and state. No redesign of the nine-game catalogue is needed.

## Offline and empty states

Cards show only verified available data. Missing content names, no eligible sessions, pending sync, deleted images and unavailable exports each need a quiet fallback. Avoid showing old suggestions as if they describe today.

# A1 / Conversation prompts

Existing C4 session data becomes a factual conversation aid

## Scope and first implementation

Generate at most two short fixed-template observations from a valid completed or deliberately stopped session with resolved activity. Prefer robust facts such as the first matched pair or the picture used. Suppress observations for interrupted, incomplete or conflicting event sequences.

Resolve opaque content IDs to caregiver-approved names on the authorized device. Missing or removed content uses a generic label or suppresses the observation. No names, photos or personal words need to enter game telemetry or an external language model.

## Correction to the additions document

It assumes universal item\_presented-to-resolution intervals already exist. The current guide defines game-specific events, such as card\_revealed, pair\_resolved and word\_selection. A shared content-use event contract or explicit per-game adapter is required before a universal observation feature can work.

A picture visible beside several others does not have an attributable viewing duration. If duration is reported, define actionable exposure, subtract paused/background intervals and use only contexts that support attribution. Longer duration does not establish enjoyment, familiarity or recall.

## Owner and dependency

Proposed package S41. Maharshitha owns the C4 component, Shanks frontend integration, Ruthika deterministic derivation under Pranav. Depends on shared content-event definitions and the session/metric foundation. Start with Reveal Match and suppress unsupported games until their adapters exist.

## Acceptance cases

The same valid synthetic session generates the same maximum two facts. Interruptions/gaps generate none. Replayed uploads do not duplicate cards. Content deletion removes sensitive display labels. Pause time cannot increase an observation duration. No template describes emotional state or cognitive ability.

## Pitch explanation

Describe the feature as helping caregivers begin a conversation using recent app activity. Claims that it improves communication, wellbeing or memory would require separate evidence.

# A2 / Played-together mode

Existing handover and game wording gain companion context

## Interface and behavior

C5 offers a toggle before starting. Shared P3/P4/P6 strings adapt naturally; all games receive strings through the common host. C4 and doctor history/report views label the session as companion play. No game-specific scoring branch is necessary.

## Snapshot and support

Freeze companion\_mode as solo or together in the session configuration. Keep caregiver assistance and hints separate: a person sitting nearby does not tell us precisely how much help was given. If support changes during play, use the existing support\_changed design or end/restart with a new snapshot; never relabel earlier events silently.

## Reference and difficulty rules

Separate companion from solo comparison series. A patient who only plays with a companion can have a clearly labelled companion reference, but those sessions do not trigger automatic level increases. Caregiver-selected settings and explicitly reviewed changes remain possible. Unknown support never becomes unassisted by default.

## Evidence interpretation

The cited Cochrane review mainly evaluated group cognitive stimulation and also included individual approaches. A companion toggle alone does not reproduce those programmes or establish their effects for Tesseract. Present this as supported participation, not a proven treatment mechanism. [2]

## Owner and dependency

Proposed S42, owned by Shanks with Maharshitha on C5, Pranav/Ruthika on schema and eligibility. Expand S04-S06, S10/S18, S16/S17 and doctor summaries. Its acceptance includes the full pipeline, not merely changed words.

## Acceptance cases

Toggle state reaches stored snapshot; shared copy changes; session/history/report context agrees; companion sessions cannot accidentally trigger level-up; offline resume preserves mode; caregiver assistance remains independently represented.

# A3 / Content curation and rotation

Describe content exposure accurately before suggesting replacement

## Correct the usage vocabulary

The source example labels an item served five times but never chosen as used rarely. That confuses exposure with selection. Record presented count, selectable-opportunity count, explicit chosen count, skipped count and last-presented time separately. In auto-assigned games there may be no meaningful choice signal at all.

C3 can show Not presented yet, Presented 5 times, or Chosen 2 of 4 times offered. Do not imply a patient fails to recognize a person. An item never offered should first receive eligible rotation, rather than a replacement nudge blaming the caregiver's content.

## Aggregation and selection

Use opaque content IDs and content versions. Derive aggregates from validated item events and deduplicated sessions. Keep the completed-event watermark so replay does not double-count. Distinguish whole-session stopping from explicit item skipping. A cancelled action is not automatically an item rejection.

Select only caregiver-approved, cached/available content that fits game, language and accessibility constraints. Then rotate by exposure and recency. Preserve removal/skip preferences; never force unused content merely to balance counts. Offline selection needs cached metadata and a defined merge policy after reconnect.

## Comparison policy

Record exact content version plus a reviewed task-complexity category, including repeated exposure. A material task change starts a separate series. Avoid restarting the reference for every new image if the defined task remains comparable; otherwise rotation could prevent enough comparable sessions accumulating.

## Owner and dependency

Proposed S43. Ruthika owns bounded aggregates under Pranav; Maharshitha C3/C4 components and Shanks selection/cache integration. Expand S09/S29 content services and S04/S13/S17 contracts/calculators. S41 and S43 share event definitions, not inferred preference scores.

## Acceptance cases

Never-presented and never-chosen cases have different labels. Counts use correct denominators. Duplicate sync changes no totals. Removing/replacing content preserves old session provenance without retaining unauthorized labels. Rotation does not silently alter difficulty or empty the activity catalogue.

# A4 / Accessibility fitting

Prototype a measurable interface suggestion before promising automation

## What is wrong with the original measurement

A magnitude-only offset cannot tell whether taps land below a button. In ordinary games the intended target may be unknown, and an off-centre successful tap is not an error. G1/G6/G8/G9 also do not currently share an `attempt_resolved` event. The feature needs an explicit measurement design and mapping to each game.

## Safer first version

Keep caregiver-adjustable text, target size and labels available immediately. Offer an optional brief tutorial with a known requested target. Process tap geometry locally. Capture successful target hits and carefully defined near misses, target size and layout/configuration version. Do not transmit raw pointer trails for this feature.

If only distance magnitude is stored, use nondirectional wording such as Try larger targets. Any directional statement requires a locally computed directional summary and enough valid observations. The known-target tutorial supports attribution; arbitrary game taps do not automatically do so.

## Eligibility and application

Two sessions alone is not a sufficient sample rule. Define minimum eligible trials, repeatability across sessions, allowable layouts and a versioned threshold after prototype testing. Keep these interface-fit values separate from performance metrics and level rules. Always require caregiver Apply, with Not now and Undo/reset. Do not auto-apply after repeated acceptance.

Record text scale, target dimensions, labels setting, input mode and `accessibility_config_version` explicitly. The current snapshot does not enumerate all of these yet. Accepted changes create a new comparison context. Unused reading games do not prove reading or vision difficulties and must not automatically disable activities.

## Owner and dependency

Proposed S44. Pranav owns the contract and suggestion policy; Shanks the shared tutorial/instrumentation; Maharshitha C7 and the suggestion card; Kovid usability cases. Define before G1/G6/G8/G9 contracts freeze, or use optional versioned fields that older sessions can omit.

## Acceptance cases

Valid off-centre hits alone produce no suggestion. Missing/unknown targets produce none. Enough eligible synthetic near misses can produce a reviewed target-size suggestion. Apply/Undo versions settings. Interface-fit fields cannot enter level-up inputs. Directional copy is impossible from magnitude-only data.

# A5 and A6 / Routine context

Optional activity bridges and caregiver daypart preferences

## A5 reminder bridge

C6 adds an offer-short-activity setting, default off. After acknowledgement, an eligible reminder may offer a skippable short activity using the normal P3/P4/P6 path. The host freezes a short-session configuration, `origin=reminder_bridge` and approved content. No session is created merely because the offer was displayed.

Keep `bridge_offered` and `bridge_accepted` as separate bridge events linked by occurrence ID; do not replace the reminder's acknowledgement state. Decline is not a failed game or disengagement signal. No recorded refusal is needed. Return to the previous safe context and never start a second game over a paused session.

Short sessions form a separate reference context from normal sessions. The reason is actual duration/task configuration, not just which button launched them. Acknowledged means a button was pressed; replace the source document's phrase reminder satisfied with reminder acknowledged.

**Proposed S45:** Shanks owns lifecycle/scheduler integration, Maharshitha C6/card controls and Pranav record contracts. Depend on reliable reminders and one eligible integrated game. Test offer/decline, retry, no connection, duplicate reminder occurrence and interrupted existing play.

## A6 daypart preferences

C7 offers optional morning/afternoon/evening preferences; default is no preference. Resolve P1/P7 suggestions against approved content, accessibility, availability and patient choice. Preferences are caregiver settings, not inferred patient behavior. A fallback is necessary if no preferred content is available.

Record `local_daypart`, local time-zone context and preference/configuration version at session start. Define daypart boundaries and behavior when the clock/time zone changes. Cache the policy for offline use. Annotate D6 first; decide whether to stratify only when the comparison needs it.

The source introduction says A6 must split the comparison key, while A6 itself leaves this open. This edition resolves that conflict with record-and-annotate as the proposed default. Do not claim to detect or treat late-day agitation.

**Proposed S46:** Maharshitha owns settings, Shanks activity selection/cache and Pranav policy persistence. Optional after core reliability. Test each boundary, no preference, no matching content, offline use and time-zone changes.



# A7 / Appointment one-pager

New C8 preview; a deterministic caregiver export

## Purpose and fields

The caregiver chooses a bounded date range and included fields. The page summarizes session dates/games and counts, current approved activity/settings, selected caregiver observations and optional known clinical context exactly as entered. Include support/session-type context, generated date, last successful sync and a plain app-activity-only explanation.

## Preview and export

C4 opens proposed C8 Appointment Preview. The exact rendered page is previewed before an explicit share or print action. The caregiver chooses the destination through the device flow; the app does not automatically email, schedule or send it to the doctor portal.

Use a deterministic template, independent of the optional doctor-side language generator. Stable input and template version must produce the same visible content/layout. Byte-identical PDFs require controlling timestamps and identifiers too; visual reproducibility is the practical acceptance criterion unless byte identity is explicitly required.

## One-page constraint

Long histories and notes will not all fit. Use a bounded range and aggregate session counts; show how many records are summarized. Let caregivers shorten/select observations and clearly label omitted detail. Do not silently clip text or shrink it until unreadable. An empty range should produce an honest no-activity summary or a prompt to change the range.

## Privacy and offline behavior

Patient-identifying exports are private even when source telemetry uses opaque IDs. Verify caregiver access at generation and export. Allow excluding clinical fields. If offline, label cached coverage and last sync; do not imply the report includes unsynced remote data. Define temporary-file storage/cleanup and make clear that sharing creates a copy outside app control.

## Relationship to D7

D7 is an assigned-doctor generated draft with a review workflow. A7 is a caregiver-selected factual appointment record. Both reuse verified metrics and provenance, but permissions, output purpose and review state differ. A7 does not bypass doctor assignments.

## Owner, dependency and acceptance

Proposed S47. Shanks owns export/share integration, Maharshitha preview/layout and Pranav authorized snapshot endpoint with Ruthika on summary fixtures. Depend on reliable summaries, profile data and approved activities. Test long/empty notes, selected fields, stale cache, identical visible rendering, demo labels, unauthorized access and cancellation before share.

# Corrected shared contracts

More than one extra field: the common host and analytics need coordination

| Contract area          | Proposed addition / rule                                                                                                                                             |
|------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Session context        | companion_mode, origin, short/normal task configuration, local_daypart, time-zone context and all relevant version identifiers. Older unknown values remain unknown. |
| Accessibility snapshot | Explicit text scale, target size/layout, labels/support, input mode and accessibility_config_version. Do not assume these already exist.                             |
| Content-use events     | Define presentation, selectable opportunity, explicit choice, resolution and skip. Map each supported game to them or declare the signal unavailable.                |
| Timing                 | Actionable foreground intervals only; subtract pauses; simultaneous visible items do not imply separately attributable viewing time.                                 |
| Interface-fit events   | Known-target tutorial first; local geometry processing; bounded derived values with target/layout metadata. Direction cannot come from magnitude alone.              |
| Content aggregates     | Opaque item/version IDs, presentation/choice-opportunity/selection/skip counts, last presentation and processed-event watermark.                                     |
| Reminder bridge        | Occurrence-linked offered/accepted events alongside existing acknowledgement; no decline-as-failure metric.                                                          |
| Exports                | Authorized selected date range/fields, source coverage, template version, exact preview and explicit destination action.                                             |

## Migration and replay

Version schemas before game owners consume them. Optional fields allow older records to remain readable, but unknown is not solo/unassisted. Backfill only from facts already present. Deduplication, content deletion and event-sequence validation apply to every new aggregate and derived card.

## Game owner handoff

Shanks provides the shared adapter and example. Aryan, Ruthika and Pranav add per-game content mapping, short-config eligibility where applicable and fixtures. Some games cannot support every addition initially; hide unavailable observations instead of inventing data.

# Comparability and AI boundaries

Preserve interpretable evidence as context becomes richer

## Separate contexts when the task actually differs

Companion and solo series stay separate. Assisted/unknown sessions cannot become unassisted references. Normal and short reminder activities use distinct configurations. Game version, input mode, support, accessibility and material task complexity remain part of comparison eligibility.

## Record without unnecessarily fragmenting every chart

Home versus Choose Activity launch origin does not necessarily alter the task; record it without automatically creating separate references. Record daypart/time zone and annotate initially. Store exact content versions and reuse exposure, but define material complexity changes rather than splitting on every rotated picture.

## Provisional reference and suggestions

Retain the proposed three comparable eligible sessions as an engineering reference, not clinical validation. Label which mode/configuration supplied it. Insufficient samples produce a hold/no-comparison state. Companion references can describe companion play while remaining ineligible for automatic difficulty increases.

## Three independent outputs

**Game performance:** existing per-game metrics and reviewed difficulty. **Content usage:** exposure, choice opportunities and rotation. **Interface fit:** known-target interaction observations and caregiver-approved accessibility settings. Do not blend these into one patient score.

## Human control

Accessibility suggestions, content removal and daypart preferences stay caregiver-controlled. A7 exports follow explicit preview/share action. Conversation cards and usage labels use fixed factual templates. Existing doctor draft generation cannot invent measurements or override reviewed settings.

## Why this matters

A higher score after larger controls, help from a companion or a much shorter activity can reflect the changed conditions. The app must expose those conditions rather than imply recovery or decline. These additions increase practical support; they do not establish new clinical capabilities.

# Build packages for the additions

Proposed S41-S47; preserve all existing S01-S40 IDs

| Package / lead          | Dependencies                               | Completion evidence                                                                                |
|-------------------------|--------------------------------------------|----------------------------------------------------------------------------------------------------|
| S41 / A1<br>Maharshitha | Revised S04/S05; S12/S13/S15               | C4 fixed facts match validated events; gaps suppress cards; privacy-safe labels.                   |
| S42 / A2<br>Shanks      | Revised S04-S06; S10/S16-S18               | Companion context persists offline, changes wording and separates analytics/level eligibility.     |
| S43 / A3<br>Ruthika     | Revised S04/S05; S09/S13/S29               | Correct count denominators; replay-safe aggregates; eligible rotation and content deletion.        |
| S44 / A4<br>Pranav      | Measurement design in S04/S05; S06/S14/S23 | Known-target prototype and minimum-sample rule; caregiver Apply/Undo; no metric contamination.     |
| S45 / A5<br>Shanks      | S18/S20/S25; revised snapshot              | Reminder acknowledged independently; optional short session isolated; no double-start after retry. |
| S46 / A6<br>Maharshitha | S09/S18/S25/S29                            | Preferences, clock boundaries, cache and fallback work; trends annotated.                          |
| S47 / A7<br>Shanks      | S09/S15/S17/S18; export contract           | Exact selected preview, readable one page, correct freshness and explicit export.                  |

## Ownership and status

Leads above are proposed integration owners, supported by the handoffs on each addition page. All packages are planning-only and NOT STARTED. No new calendar estimate is committed. Existing five/two/two game counts remain unchanged.

## When these become release dependencies

If all seven are accepted into the full release, S39 gains S41-S47 as prerequisites and S40 still follows S39 and S30. If A4/A6 remain optional experiments, list that explicitly in the release scope. Do not create a cycle by making an addition depend on S39 while S39 depends on it.

## Current source-of-truth distinction

The current canonical roadmap remains the 40-package revision. This PDF presents the proposed extension with reviewed corrections. Recording a new PDF is not equivalent to accepting each experimental threshold or completing any feature.

# Existing build plan / required updates

Integrate the additions into existing packages, then verify the full release

| Existing packages | What the next revision must include                                                                                       |
|-------------------|---------------------------------------------------------------------------------------------------------------------------|
| S01-S03           | Record adopted additions, priorities, exact owners and C8 versus preview-subview choice; update Figma.                    |
| S04-S06           | Freeze new context/content/interface contracts; build one shared host example and migration fixtures.                     |
| S09 / S29         | Content metadata, rotation, authorization, versioning and cached selection rules.                                         |
| S10 / S15 / S18   | Companion handover, factual caregiver cards and context-safe patient return paths.                                        |
| S12-S14           | Ingest new versions, dedupe aggregates and preserve mode/settings across offline replay.                                  |
| S16 / S17         | Eligibility for companion/short/support contexts; reference and trend annotations; no accessibility input in level rules. |
| S20 / S25         | Optional bridge, independent reminders, local schedules, duplicate-occurrence and restart behavior.                       |
| S23 / S36-S38     | Usability and doctor-view consistency; report provenance and patient assignment checks.                                   |
| S39 / S30 / S40   | End-to-end addition cases; truthful evidence-based pitch; release only the agreed tested scope.                           |

## Recommended delivery sequence

First design the common context/content contracts and A4 measurement option. Add A2 and a narrow A1 with the integrated loop. Once reliable summaries/content exist, implement A7 and A3. Add A5 only after independent reminders pass. A6 remains a simple optional schedule; A4 automation follows prototype evidence.

## Effort and workload

Estimate design, mobile, backend, migration and QA separately after a shared example exists. Avoid assuming a toggle or one-page export takes only its UI coding time. Pranav and Shanks need protected integration windows; Ruthika's shared analytics workload must not collide invisibly with her two games.

## Documentation handoff

A future accepted plan revision must synchronize the product/UI guide, handbook, roadmap and master context, then regenerate their PDFs. The original additions document and supplied PPT remain preserved as source material.

# Acceptance, sources and current status

A new planning document; no application implementation claim

## Cross-feature verification

Use synthetic sessions for zero/gapped events, completed/stopped/interrupted play, replay twice, deleted content, missing labels, companion-only histories, support changes, short sessions, time-zone boundaries, accessibility Apply/Undo, declined offers and long appointment notes. Compare caregiver cards, doctor summaries and exports to the same source facts.

Confirm all nine game modules remain within the common host; added fields are versioned; old unknown context is labelled; derived facts are reproducible; private data is not sent through game telemetry; and no new feature infers preference, adherence or disease progression from unrelated signals.

## Local sources

PS003\_ADDITIONS.md supplies A1-A7. PS003\_MASTER\_CONTEXT.md and docs/PS003\_BUILD\_HANDBOOK.md / PS003\_ROADMAP.md supply the current 40-package baseline. docs/PS003\_GAME\_ASSIGNMENTS\_AND\_UI.md supplies game/screen details. docs/PS003\_PATIENT\_EXPERIENCE.md supplies required reminders and swipe coloring. The original six-slide PPT supplies the pitch content discussed in the separate existing product specification.

## Research reference

The Cochrane evidence reviewed for A2 distinguishes group and individual delivery and does not validate a Tesseract companion toggle. The same clinical-evidence limits described in the baseline remain. [Cochrane cognitive stimulation review](#). Other baseline references remain in the separate product specification.

## Decisions still open

Exact addition priorities/owners; whether A4 automation is included or prototype-only; whether A6 is a release requirement; C8 standalone versus subview; per-game content-event coverage; sample thresholds; new schedule; and real-data retention/export policy. Defaults proposed here support implementation planning without presenting them as validated clinical rules.

## Delivery location and status

This additions-only PDF is saved directly in the SIH project root, as requested. It specifies all seven additions, their corrections and changes needed to the existing plan. No game, backend, frontend or patient-data operation was implemented by creating it. The original PPT and existing output-folder PDFs remain unchanged.